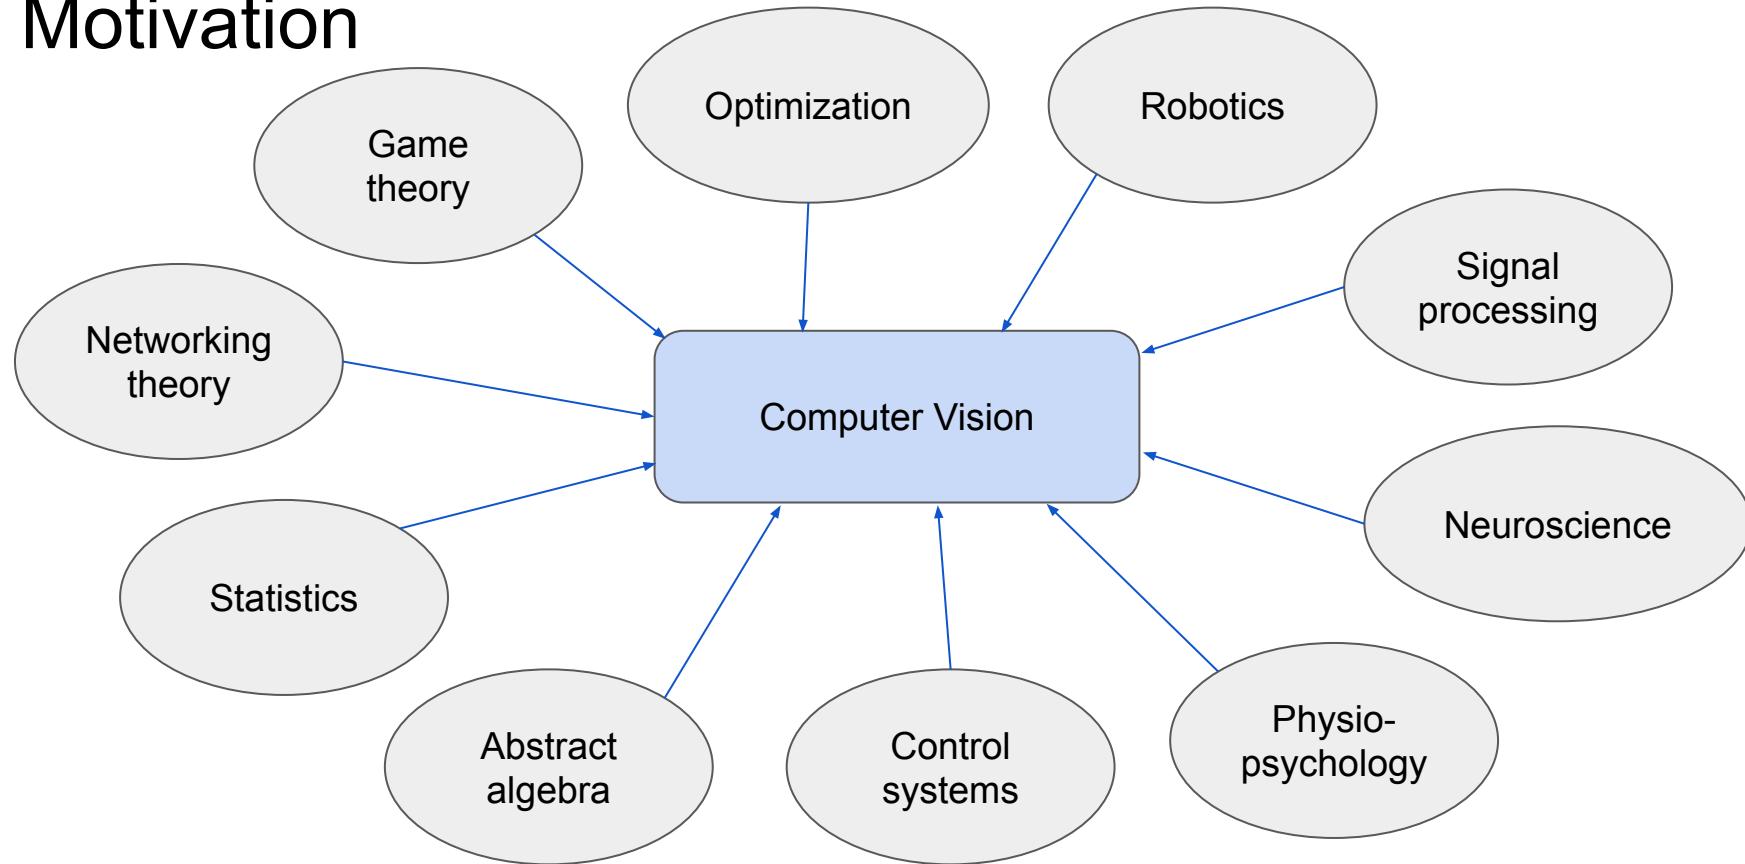


Foundations of Computer Vision

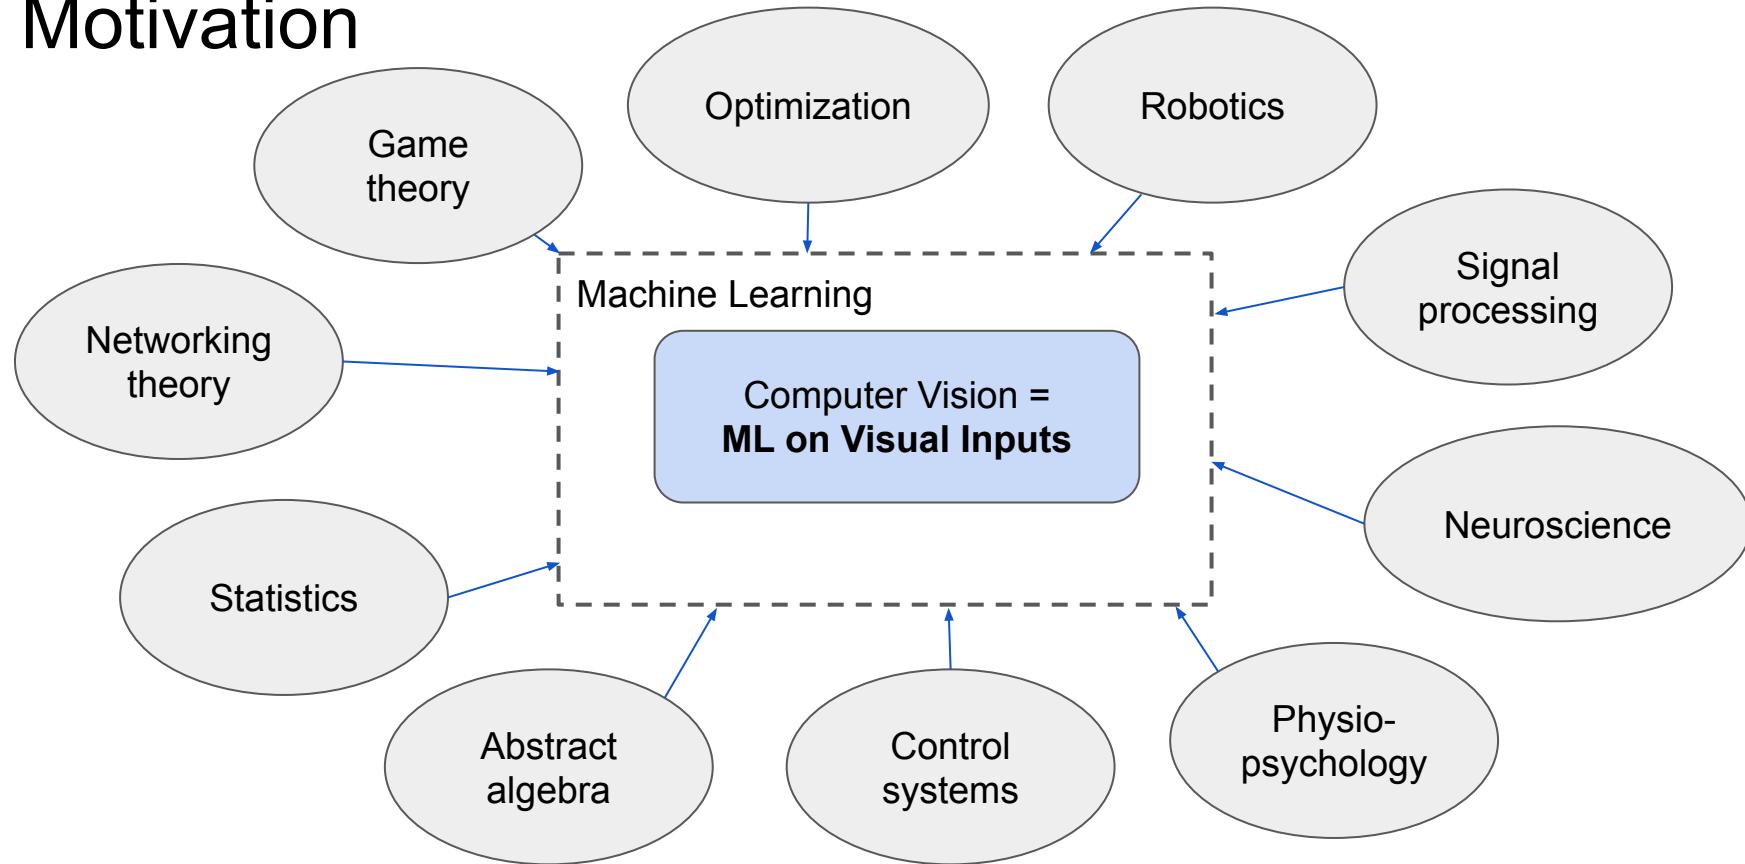
Lecture 1

Parker Koch

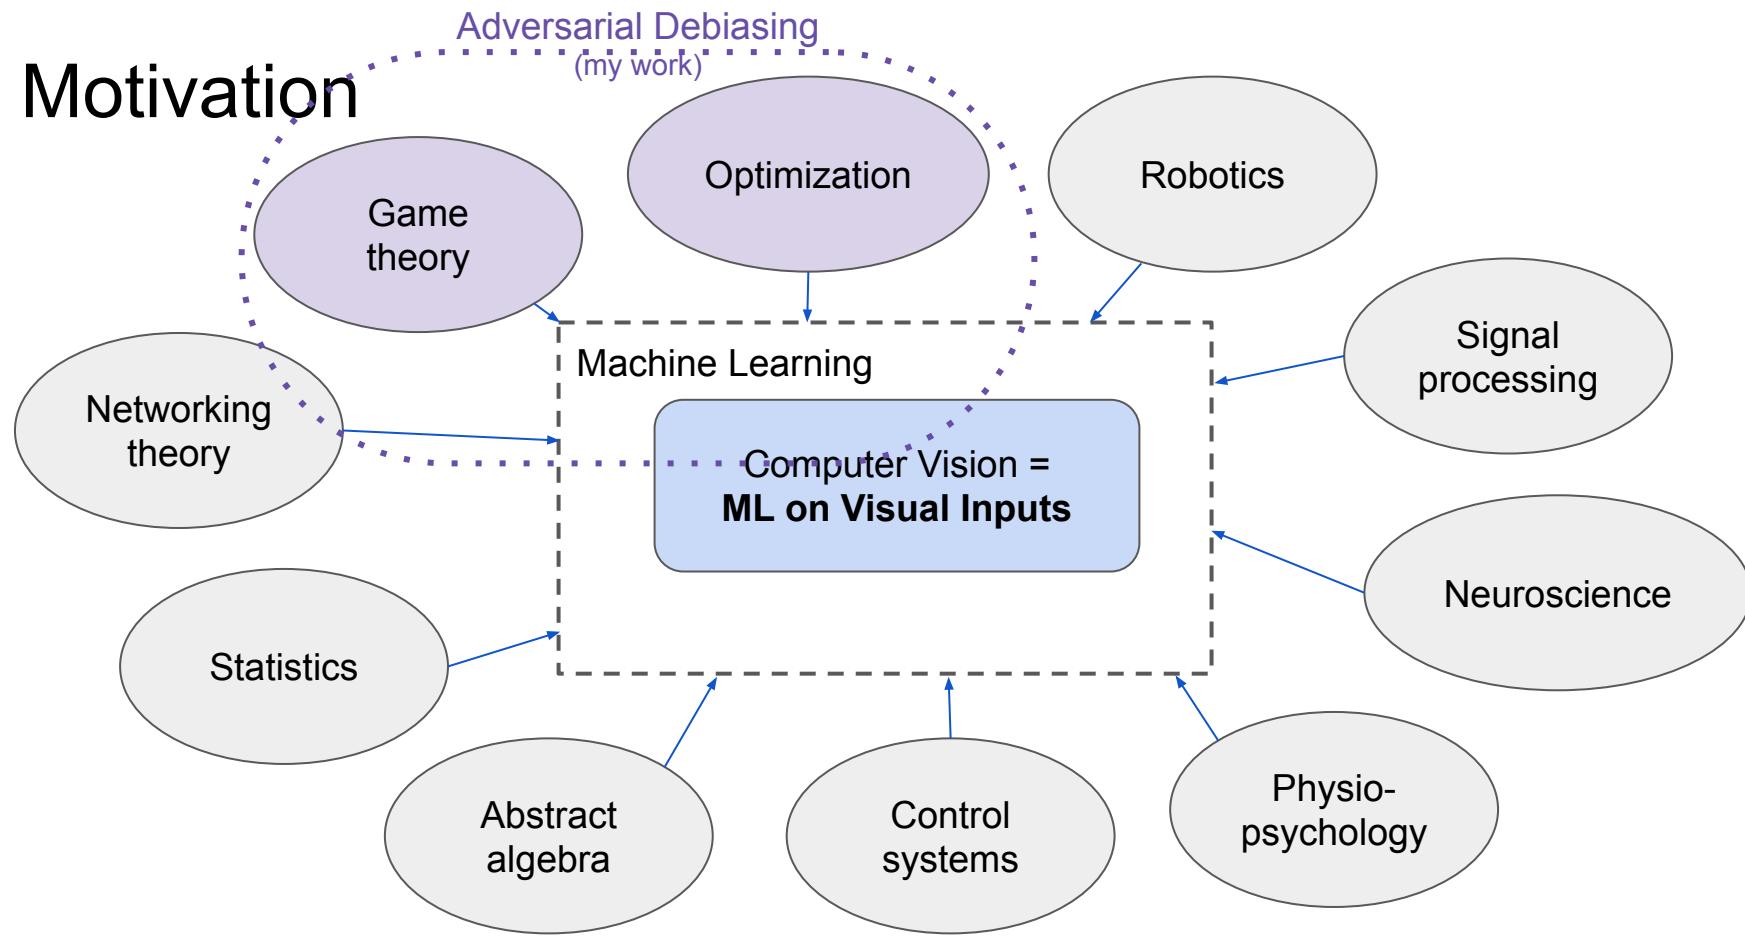
Motivation



Motivation



Motivation



Learning Objectives

By the end of this lecture, you should:

- Know how color is represented digitally
- Understand the structure of digital images and videos
- Begin thinking about images as points or functions
- Recognize the differences between kinds of image transforms
- Understand bilinear interpolation for image transformations

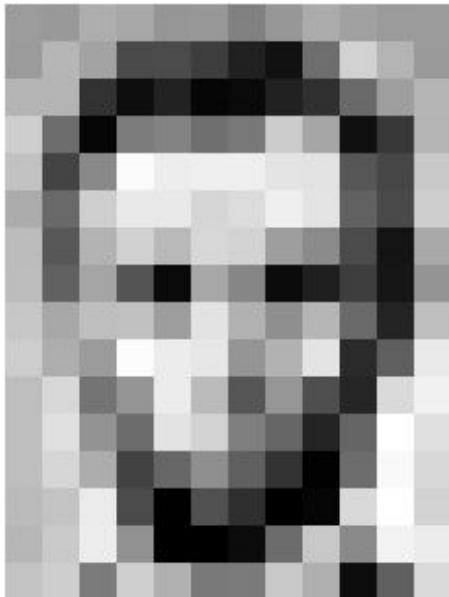
Images, Pixels, and Color

What is this an image of?

Hint: Higher numbers = whiter shades of gray

157	153	174	168	150	152	129	151	172	161	156	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	216
187	196	235	75	1	81	47	0	6	217	256	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

Images, Pixels, and Color



Abraham Lincoln!

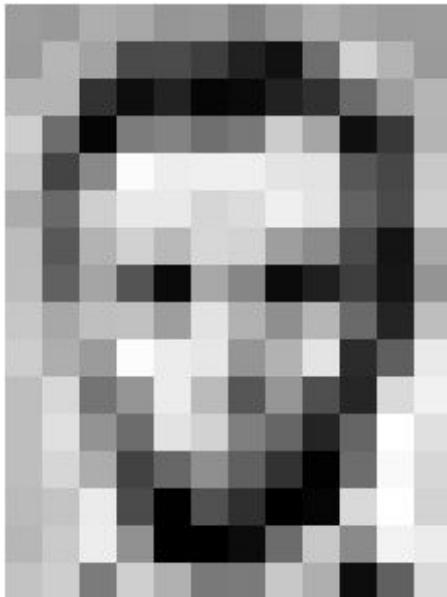
← Same information →
← Easy Hard →

Why?

“Semantic gap”

157	153	174	168	150	152	129	151	172	161	156	156
156	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	216
187	196	235	75	1	81	47	0	6	217	256	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

Images, Pixels, and Color



Our eyes and brains
evolved to identify
← this efficiently.

Computers did not.

**Vision tasks are hard
because we are trying to
mimic psychological
tasks.
Always remember that
“computer vision” is
inherently about human
vision.**

157	153	174	168	150	152	129	151	172	161	156	156
156	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	216
187	196	235	75	1	81	47	0	6	217	256	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

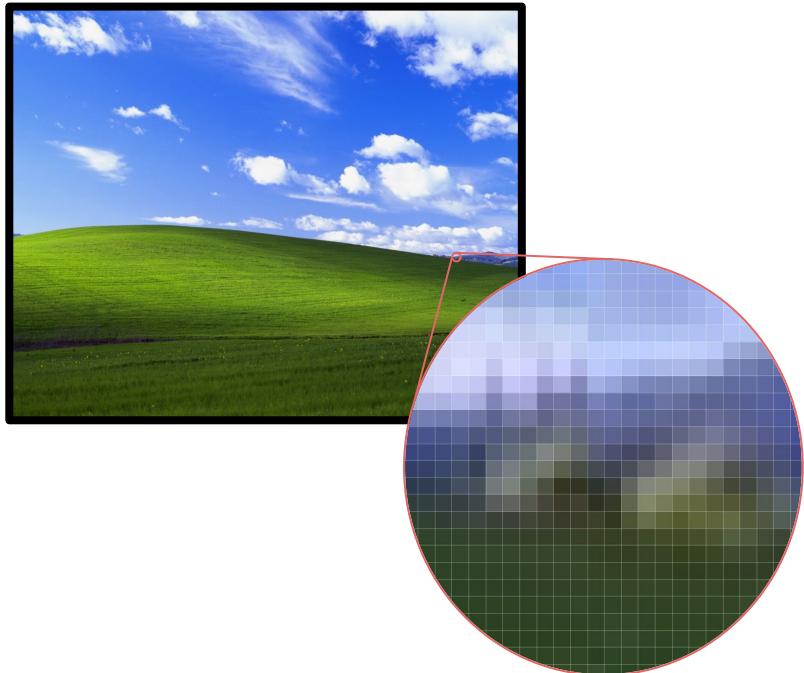
Images, Pixels, and Color

157	153	174	168	150	152	129	151	172	161	165	156
165	182	163	74	75	62	83	17	110	210	180	154
180	180	50	14	84	6	10	33	48	105	159	181
206	109	6	124	131	111	120	204	166	15	56	180
194	68	157	251	257	299	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	168	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	158	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	6	12	108	200	138	243	236
195	206	123	207	177	121	123	200	176	13	96	218

- **Digital image:** 2D array of pixels
- **Pixel:** “Picture element,” holds an atomic unit of visual data
 - Grayscale: 1D value, represents luminosity at a point

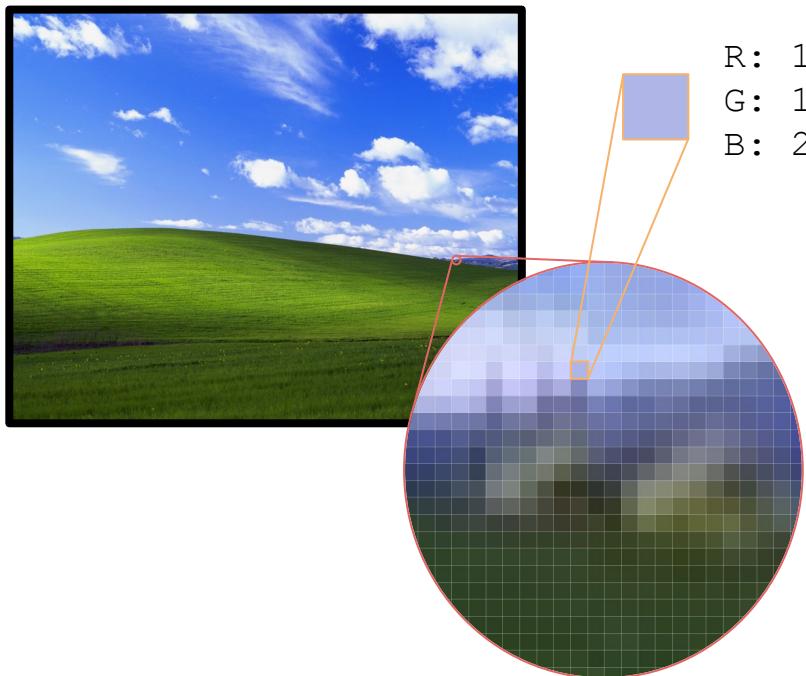


Images, Pixels, and Color

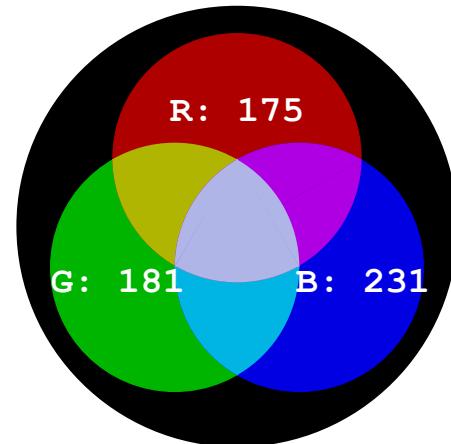


- **Digital image:** 2D array of pixels
- **Pixel:** “Picture element,” holds an atomic unit of visual data
 - Grayscale: 1D value, represents luminosity at a point
 - Color: 3D value, most commonly RGB

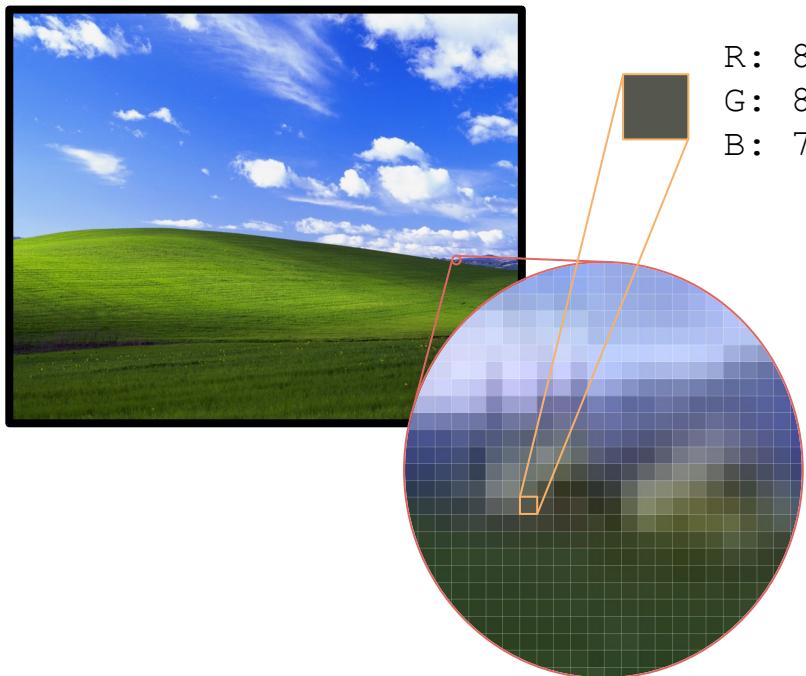
Images, Pixels, and Color



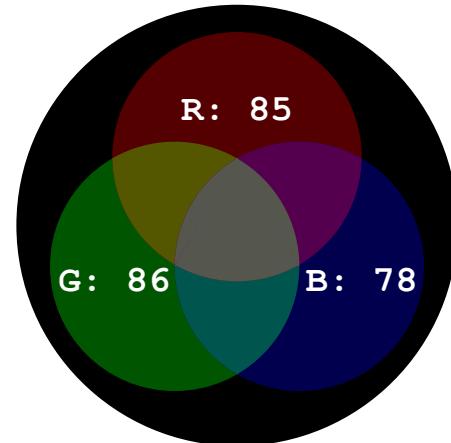
- **Digital image:** 2D array of pixels
- **Pixel:** “Picture element,” holds an atomic unit of visual data
 - Grayscale: 1D value
 - Color: 3D value



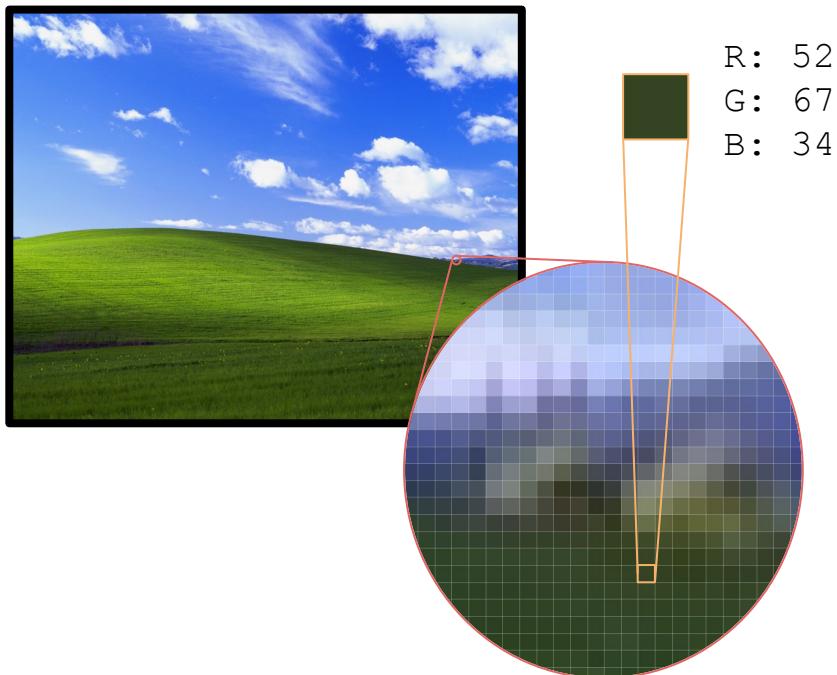
Images, Pixels, and Color



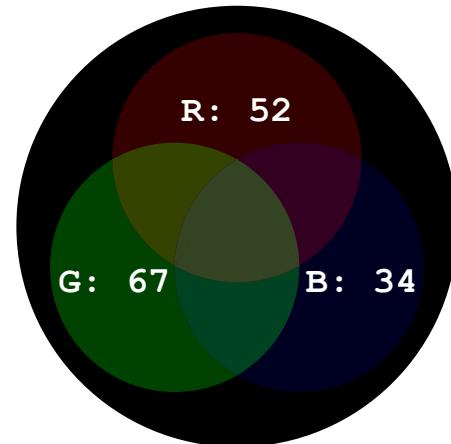
- **Digital image:** 2D array of pixels
- **Pixel:** “Picture element,” holds an atomic unit of visual data
 - Grayscale: 1D value
 - Color: 3D value



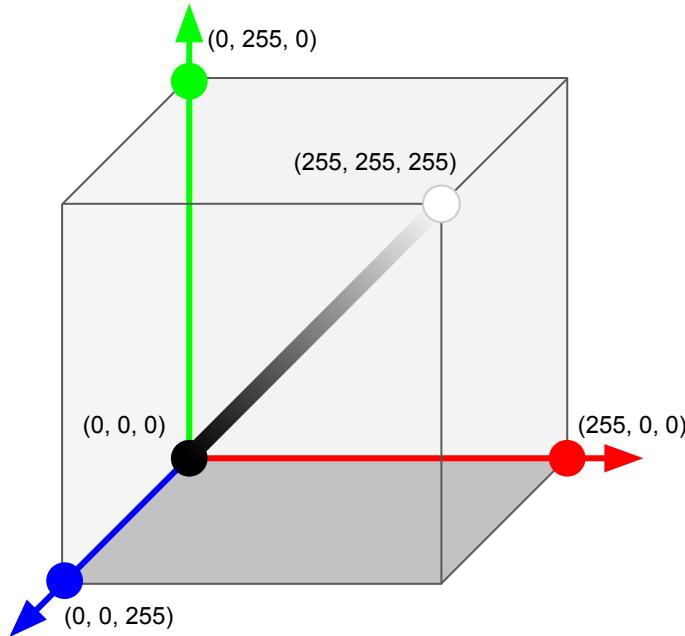
Images, Pixels, and Color



- **Digital image:** 2D array of pixels
- **Pixel:** “Picture element,” holds an atomic unit of visual data
 - Grayscale: 1D value
 - Color: 3D value

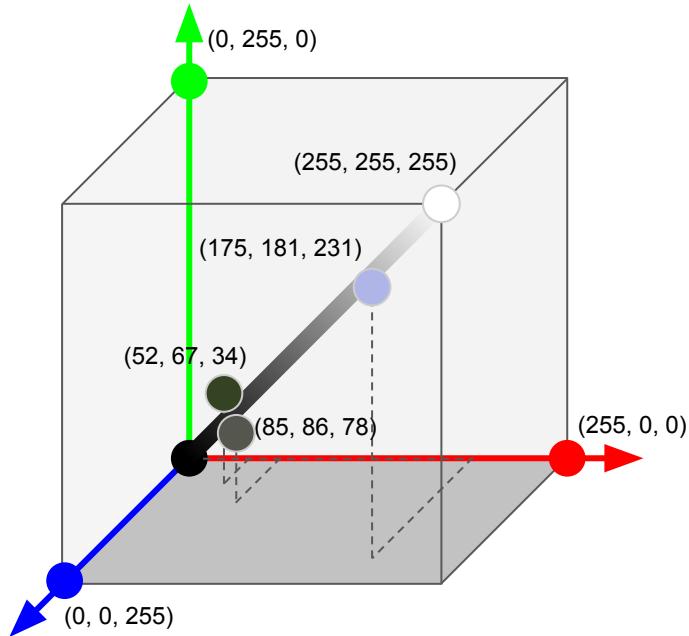


Images, Pixels, and Color



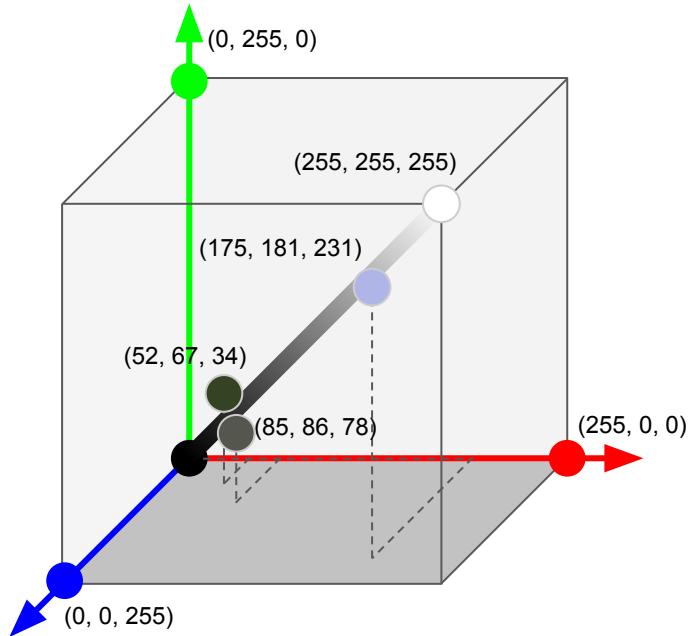
- RGB “color space”: Bounded 3D vector space
- Black = (0, 0, 0); White = (255, 255, 255)
- Gray tones = equal amounts of R, G, B

Images, Pixels, and Color



- RGB “color space”: Bounded 3D vector space
- Black = (0, 0, 0); White = (255, 255, 255)
- Gray tones = equal amounts of R, G, B

Images, Pixels, and Color



- RGB “color space”: Bounded 3D vector space
 - Black = (0, 0, 0); White = (255, 255, 255)
 - Gray tones = equal amounts of R, G, B
-
- Other color spaces exist!
 - HSV: hue, saturation, value
 - Lab: lightness, a/b
 - CMYK: cyan, magenta, yellow, black

Images, Pixels, and Color

- Channel: A slice of an image, taking only the values from one dimension of the color space
- RGB images are three-channel; grayscale are single-channel



R

Images, Pixels, and Color

- Channel: A slice of an image, taking only the values from one dimension of the color space
- RGB images are three-channel; grayscale are single-channel



R



G



B

What's happening here?

Images, Pixels, and Color

- Channel: A slice of an image, taking only the values from one dimension of the color space
- RGB images are three-channel; grayscale are single-channel



R



G



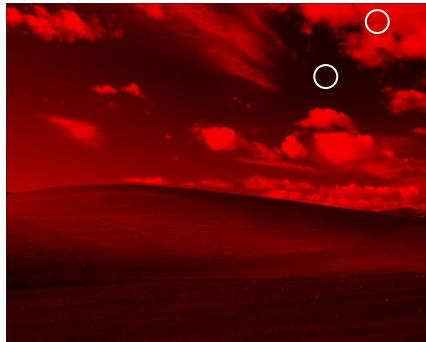
B

What's happening here?

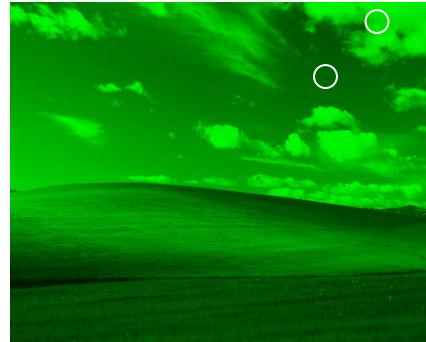


Images, Pixels, and Color

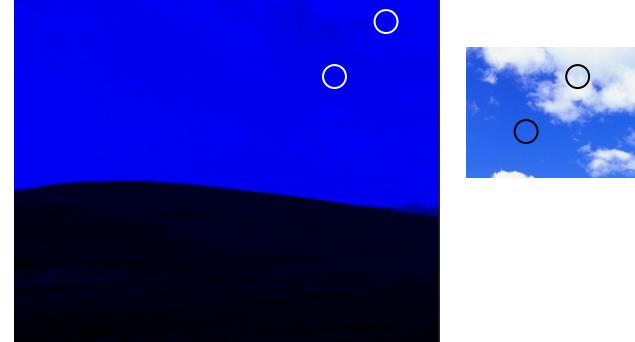
- Channel: A slice of an image, taking only the values from one dimension of the color space
- RGB images are three-channel; grayscale are single-channel



R



G



B

	Sky	Cloud	
R:	34	205	
G:	95	221	
B:	235	244	<i>Nearly equal!</i>

Video

- Video: Sequence of images (frames); 3D array of *voxels*
- Voxel: “Video pixel”
- More on video later in the course!

$t = 1 / \text{frame rate}$



Learning Objectives

By the end of this lecture, you should:

- ~~Know how color is represented digitally~~
- ~~Understand the structure of digital images and videos~~
- Begin thinking about images as points or functions
- Recognize the differences between kinds of image transforms
- Understand bilinear interpolation for image transformations

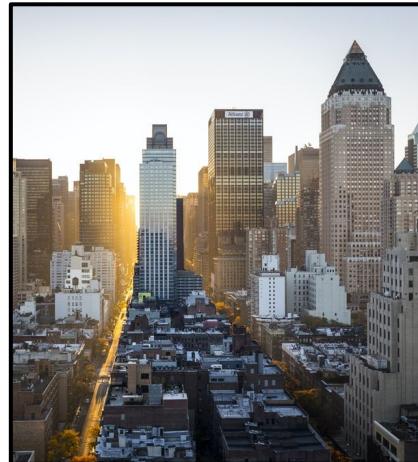
Image Representations and Transforms

- How would you take the “average” of these images?



Image Representations and Transforms

- How would you take the “average” of these images?



/2

A division by two symbol, indicating the result of the average operation.

Image Representations and Transforms

- How would you take the “average” of these ~~images~~ vectors?

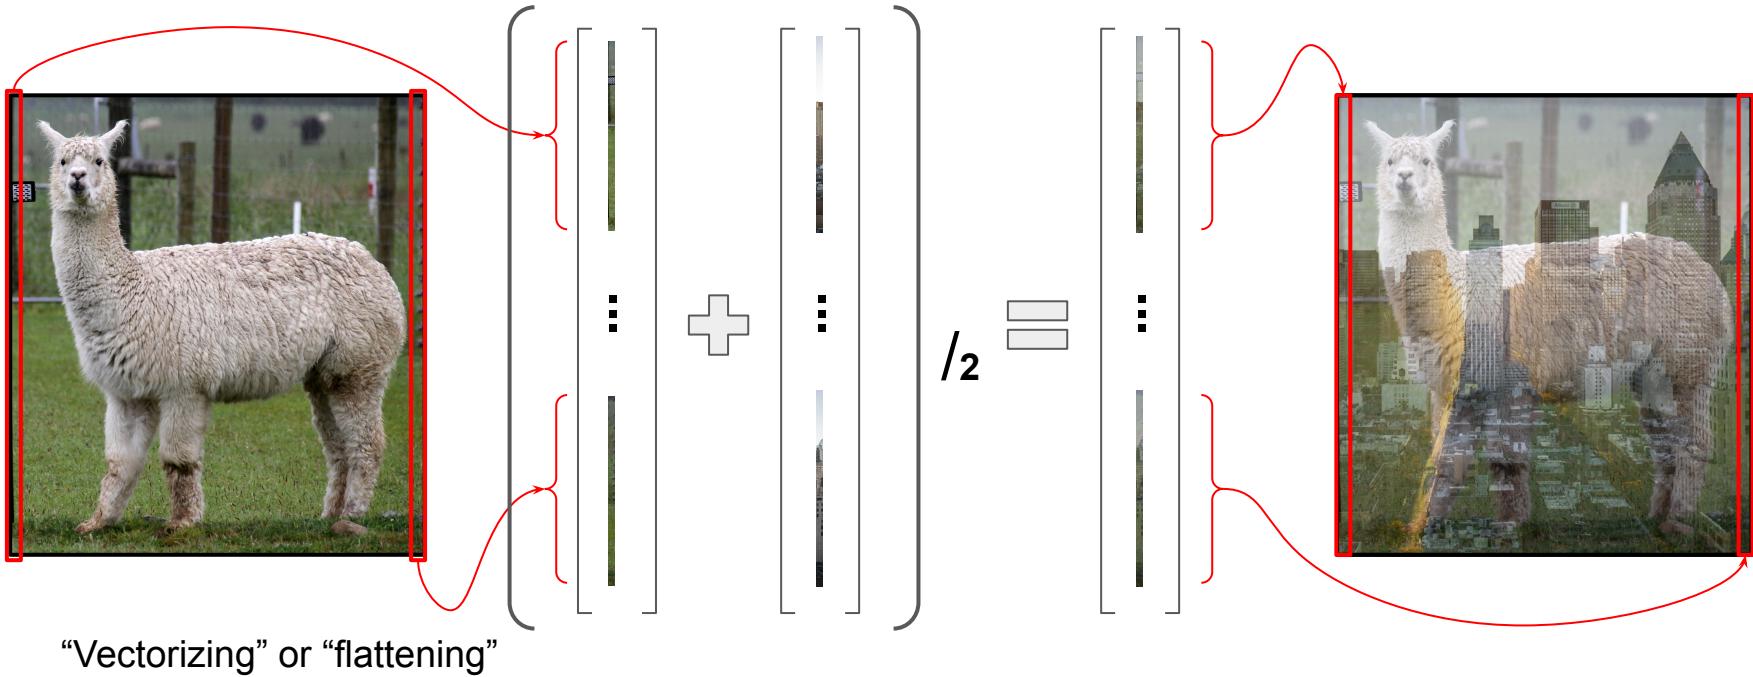


Image Space

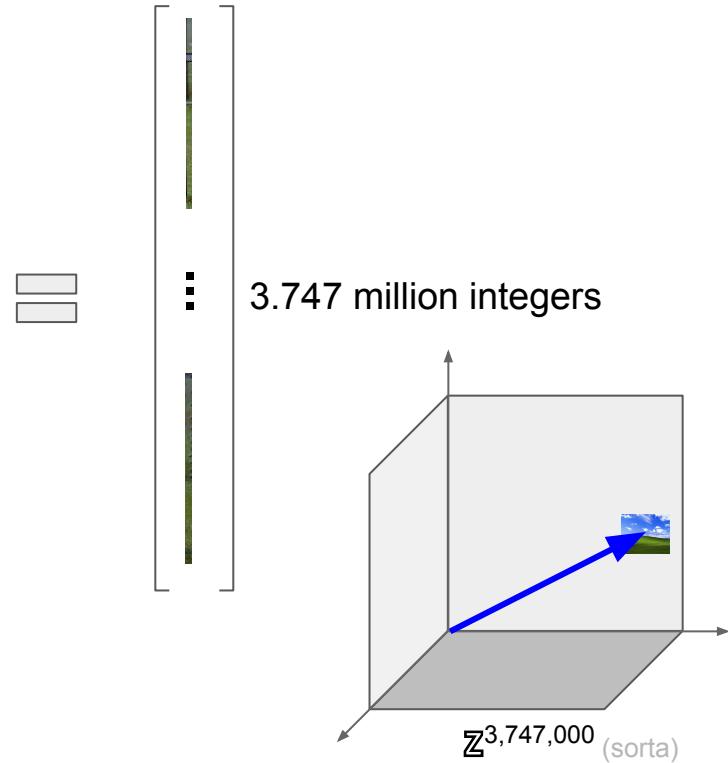
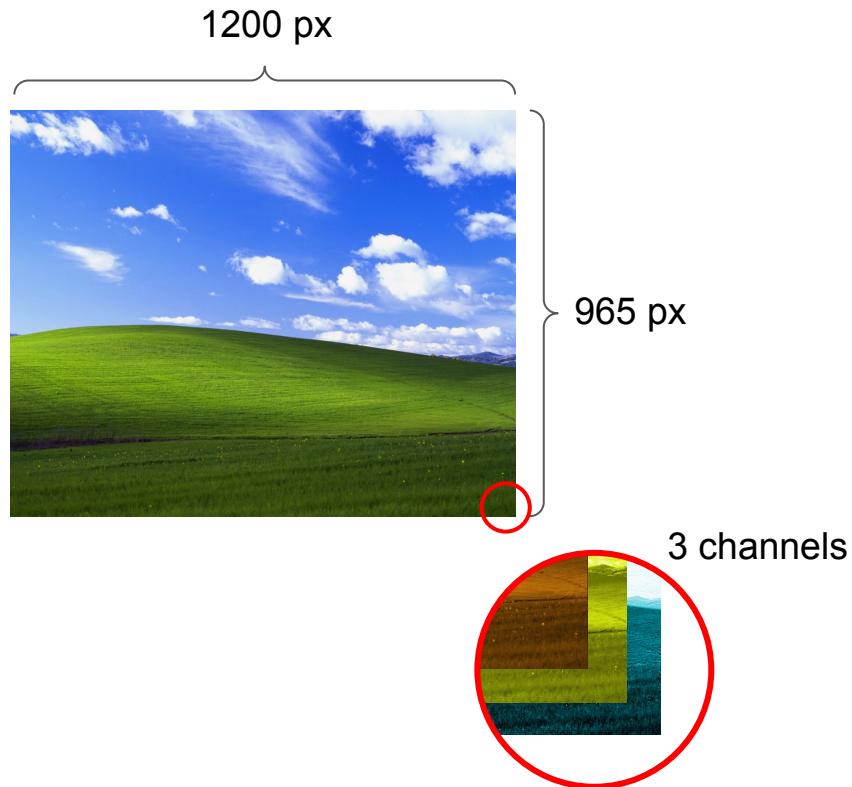


Image Space

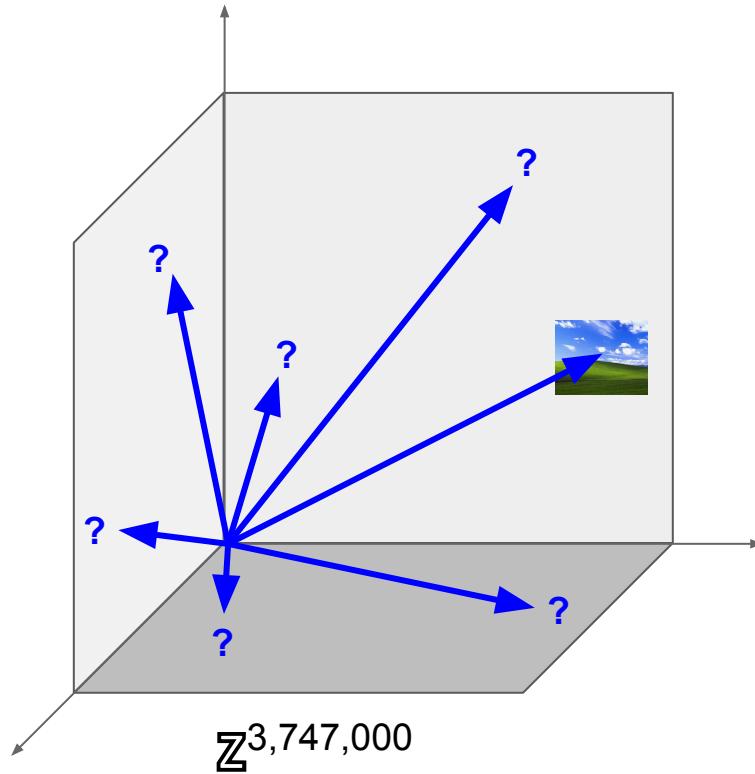


Image Space

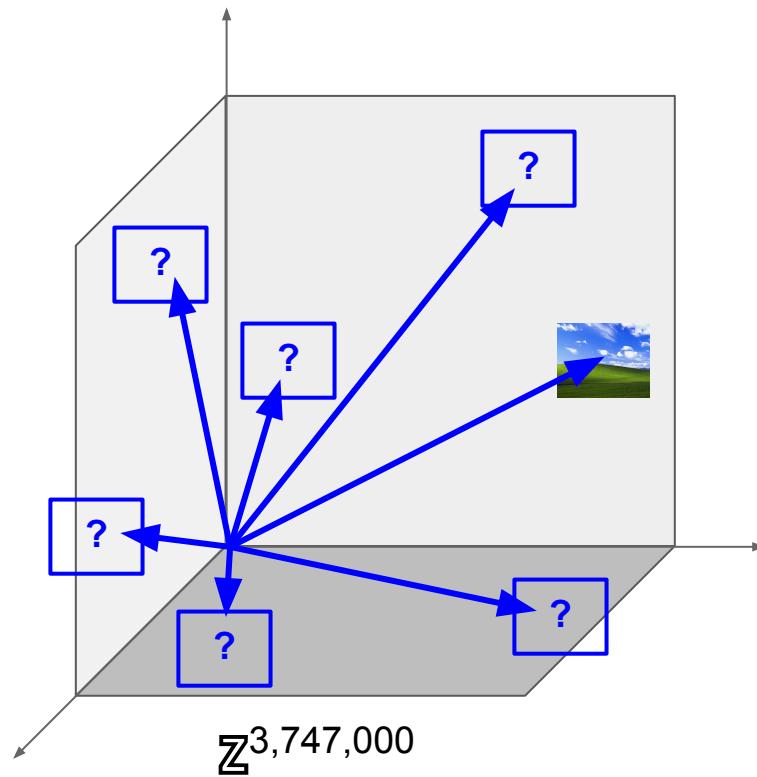


Image Space

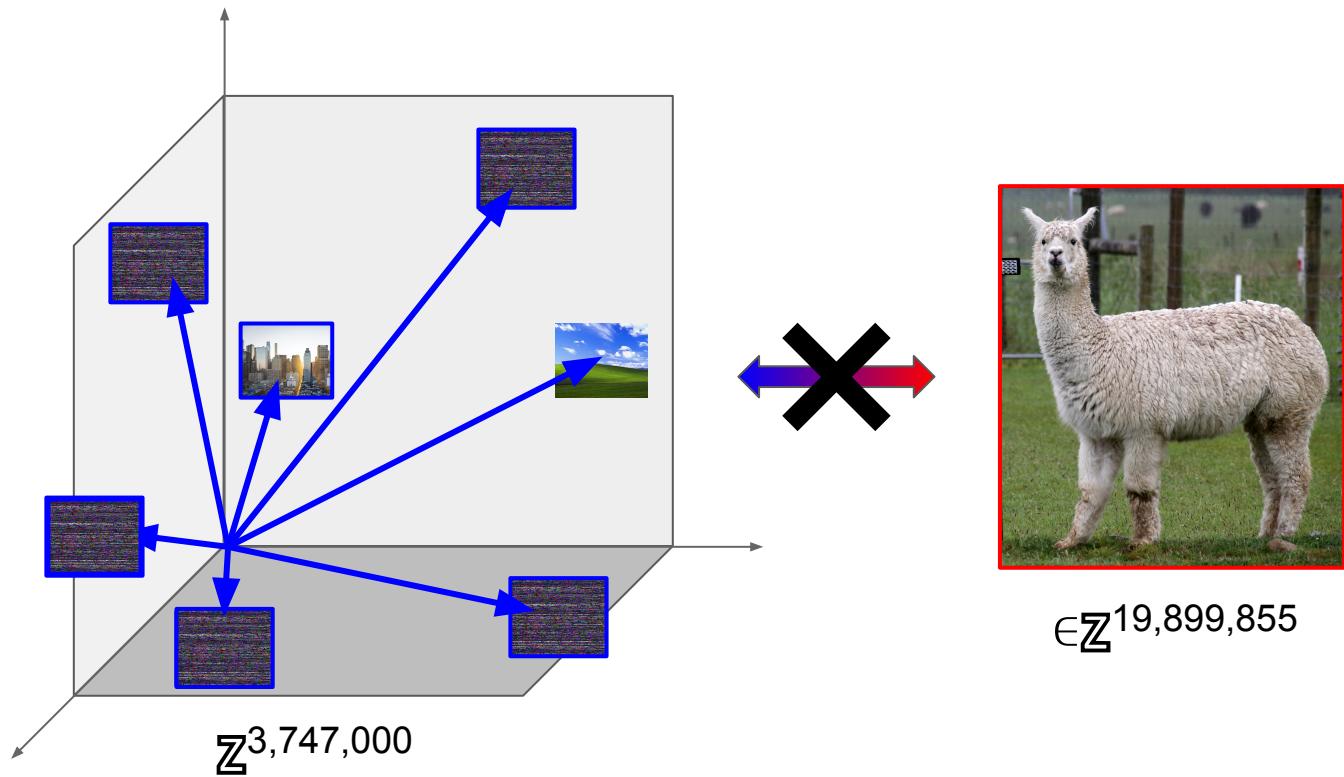
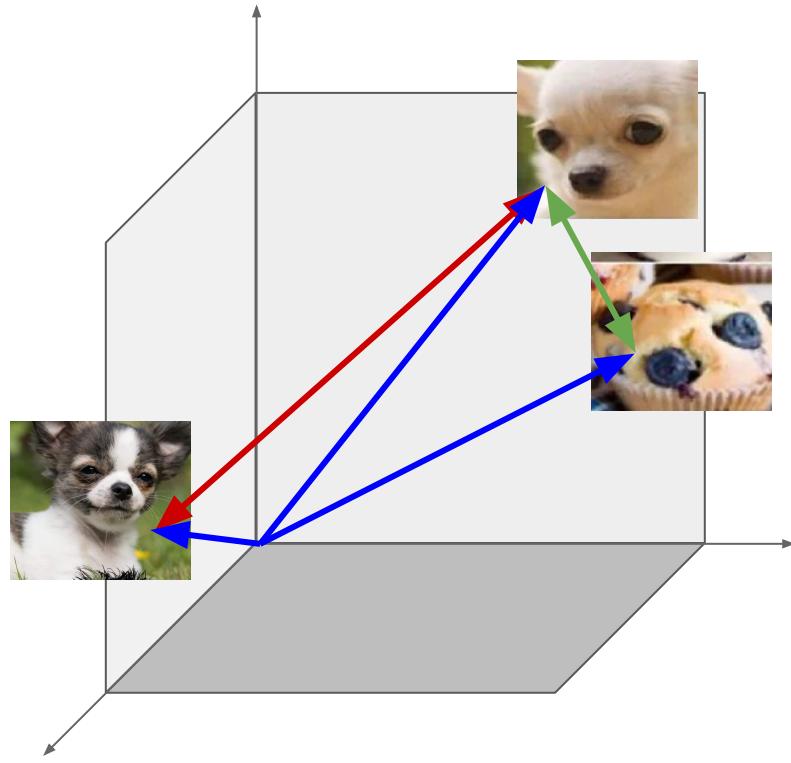
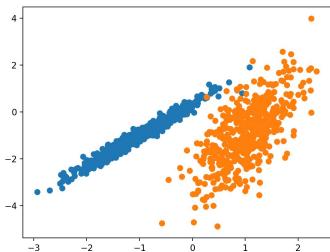
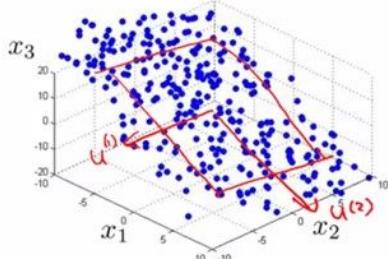


Image Vector Operations

- Distance metric
 - Pixelwise
 - Non-semantic
- PCA
- Clustering

Later in the course:

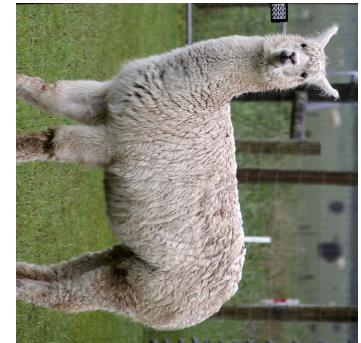


Images as Functions

Q: How do we rotate an image using its vector form?

A: Not easily.

We need to think of images as *functions* for this.



$$\begin{bmatrix} | \\ | \\ \vdots \\ | \end{bmatrix}$$



?

Images as Functions

Function: a mapping from an input space to an output space

Images map *pixel locations* to *pixel values* (colors)



Image $I(x,y)$

(0,0)	(0,1)	(0,2)	(0,3)	(0,4)
(1,0)	(1,1)	(1,2)	(1,3)	(1,4)
(2,0)	(2,1)	(2,2)	(2,3)	(2,4)
(3,0)	(3,1)	(3,2)	(3,3)	(3,4)
(4,0)	(4,1)	(4,2)	(4,3)	(4,4)

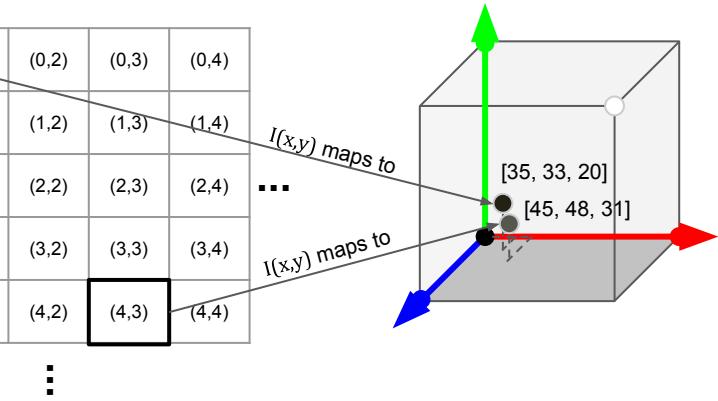
$$I(0,0) = \text{[Red Box]} = [35, 33, 20]$$

$$I(4,3) = \text{[Red Box]} = [45, 48, 31]$$

Input space: \mathbb{Z}^2

(0,0)	(0,1)	(0,2)	(0,3)	(0,4)
(1,0)	(1,1)	(1,2)	(1,3)	(1,4)
(2,0)	(2,1)	(2,2)	(2,3)	(2,4)
(3,0)	(3,1)	(3,2)	(3,3)	(3,4)
(4,0)	(4,1)	(4,2)	(4,3)	(4,4)

Output space: RGB



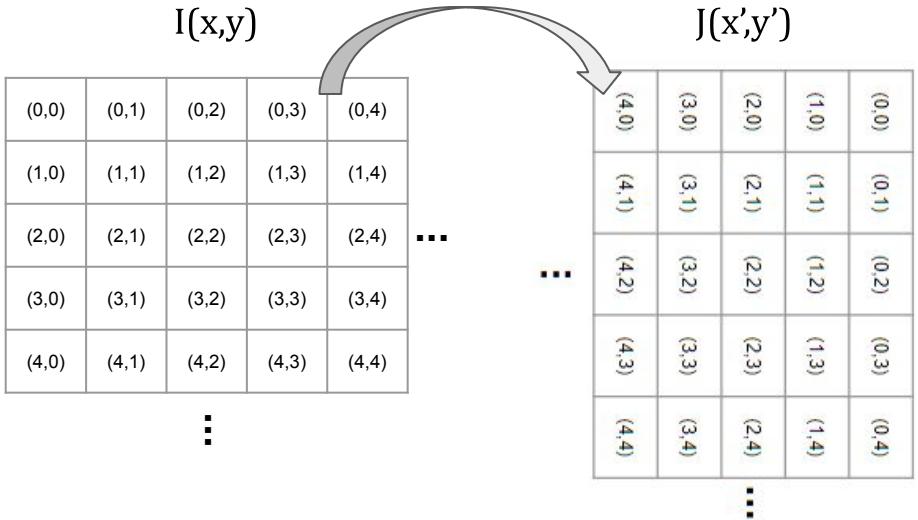
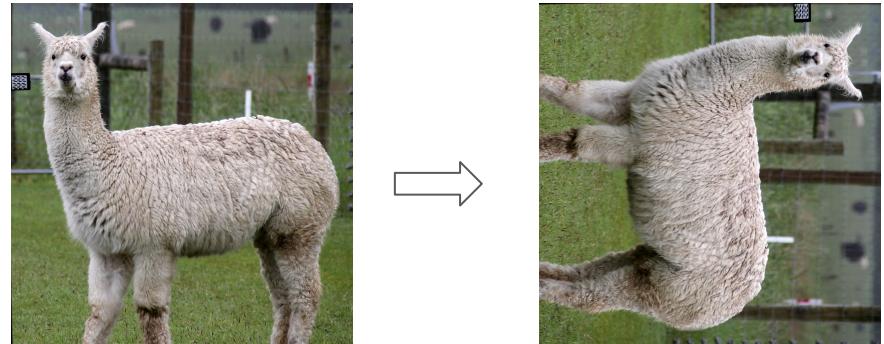
Images as Functions

Q: How do we rotate an image using its **function form**?

A: Rotate the pixel **locations**!

Domain transform: An operation on the input space of an image

- Rotation
- Translation
- Scaling (up/downsampling)
- Skew/perspective shift



Other Image Transforms

Range transforms: Operations on the pixel values, but not the locations

- Negative coloring
- Recoloring
- Brightness/contrast adjustment



Window transforms: Operations on a neighborhood of pixel values, but not locations

- Blurring
- Edge detection/highlighting

...and many that don't fit any of these categories!



Learning Objectives

By the end of this lecture, you should:

- ~~Know how color is represented digitally~~
- ~~Understand the structure of digital images and videos~~
- ~~Begin thinking about images as points or functions~~
- ~~Recognize the differences between kinds of image transforms~~
- Understand bilinear interpolation for image transformations

Interpolation (Warmup)

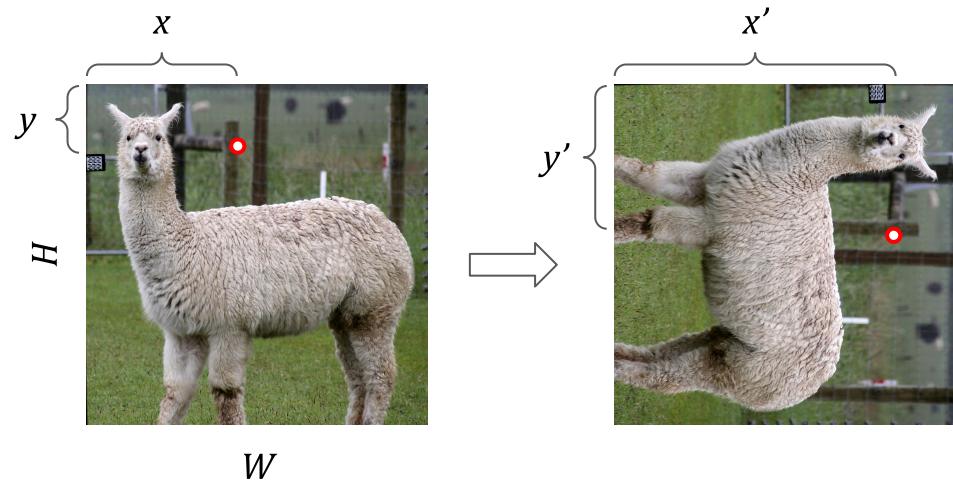
Find $J(x', y')$, the result of rotating $I(x, y)$ by 90 degrees clockwise.

Approach: Express (x, y) in terms of (x', y')

$$\begin{cases} x = & y' \\ y = & H - x' \end{cases}$$

Build $J(x', y')$ by looking up the corresponding pixels in $I(x, y)$:

$$J(x', y') = I(y', H - x')$$



General approach: Find inverse transformation, sample pixels from original image

$$J\left(\begin{matrix} x' \\ y' \end{matrix}\right) = I\left(T^{-1}\left(\begin{matrix} x' \\ y' \end{matrix}\right)\right)$$

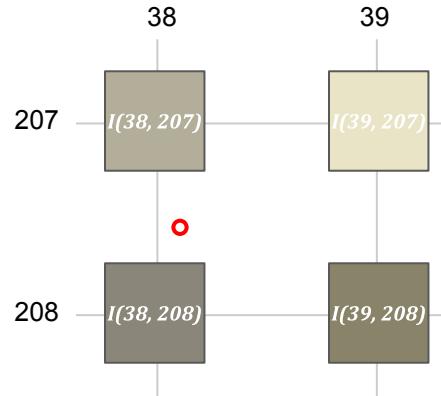
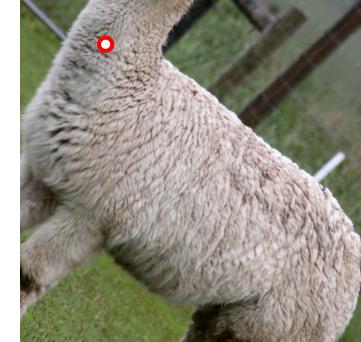
Interpolation

What if I asked you to rotate this image 45 degrees instead of 90?

What if (x,y) aren't valid, integer pixel locations?

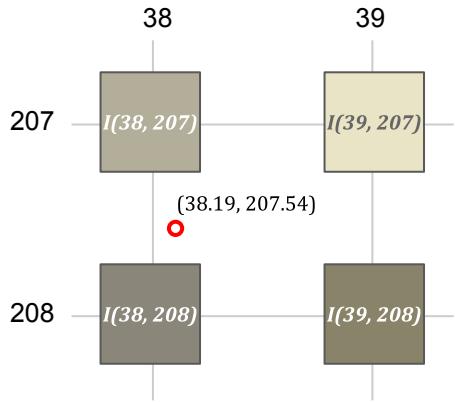
$$J(43, 17) = I(38.19, 207.54)$$

Approach: **Interpolate** between surrounding pixel values based on how close it is to each

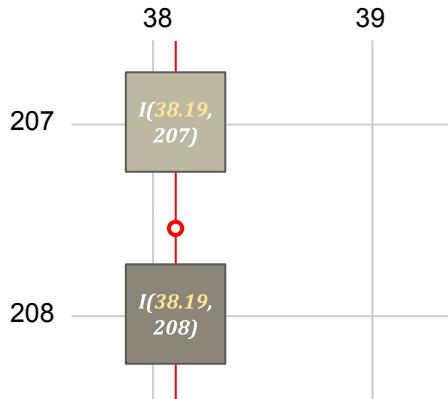


Interpolation

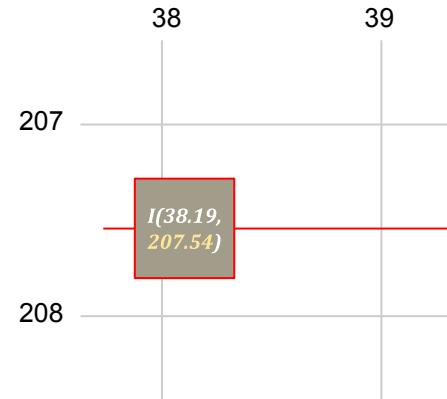
Step 1: Locate point to interpolate



Step 2: Interpolate along one axis



Step 3: Interpolate again



Learning Objectives

By the end of this lecture, you should:

- ~~Know how color is represented digitally~~
- ~~Understand the structure of digital images and videos~~
- ~~Begin thinking about images as points or functions~~
- ~~Recognize the differences between kinds of image transforms~~
- ~~Understand bilinear interpolation for image transformations~~